

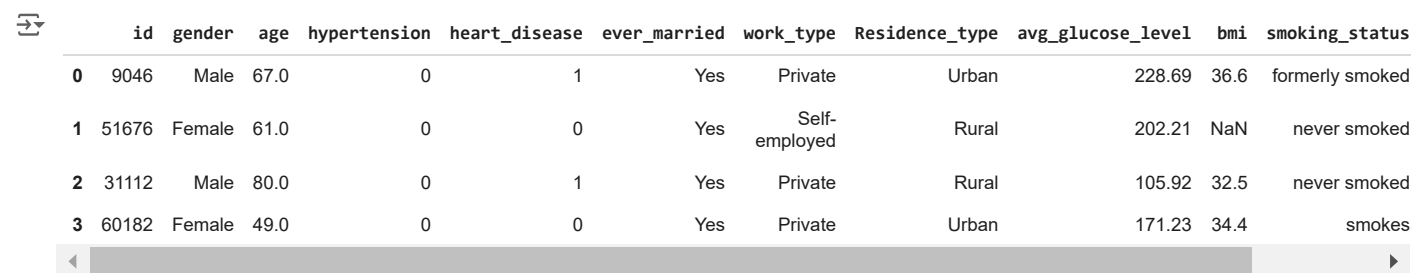
## Importing the libraries

```
import numpy as np
import pandas as pd
import seaborn as sns
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.preprocessing import StandardScaler, LabelEncoder, OneHotEncoder
from sklearn.ensemble import RandomForestClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import classification_report, confusion_matrix, roc_auc_score, roc_curve
from imblearn.over_sampling import SMOTE
import matplotlib.pyplot as plt
```

## Importing the dataset

```
df = pd.read_csv("/content/drive/MyDrive/MACHINE LEARNING PROJECTS/stroke-data.csv")
```

```
df.head()
```



	id	gender	age	hypertension	heart_disease	ever_married	work_type	Residence_type	avg_glucose_level	bmi	smoking_status
0	9046	Male	67.0	0	1	Yes	Private	Urban	228.69	36.6	formerly smoked
1	51676	Female	61.0	0	0	Yes	Self-employed	Rural	202.21	NaN	never smoked
2	31112	Male	80.0	0	1	Yes	Private	Rural	105.92	32.5	never smoked
3	60182	Female	49.0	0	0	Yes	Private	Urban	171.23	34.4	smokes

Next steps:

[Generate code with df](#)[View recommended plots](#)[New interactive sheet](#)

## Checking the structure of the data and presence of any null values

```
df.shape
```

```
(5110, 12)
```

```
df.size
```

```
61320
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5110 entries, 0 to 5109
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype  
---  -
0   id                    5110 non-null  int64  
1   gender                5110 non-null  object  
2   age                  5110 non-null  float64 
3   hypertension          5110 non-null  int64  
4   heart_disease         5110 non-null  int64  
5   ever_married          5110 non-null  object  
6   work_type             5110 non-null  object  
7   Residence_type        5110 non-null  object  
8   avg_glucose_level     5110 non-null  float64 
9   bmi                  4909 non-null  float64 
10  smoking_status        5110 non-null  object  
11  stroke                5110 non-null  int64  
dtypes: float64(3), int64(4), object(5)
memory usage: 479.2+ KB
```

```
df.isnull().sum()
```


	0
id	0
gender	0
age	0
hypertension	0
heart_disease	0
ever_married	0
work_type	0
Residence_type	0
avg_glucose_level	0
bmi	201
smoking_status	0
stroke	0

Treating the null values

```
df['bmi'].unique()
```

```
array([36.6, nan, 32.5, 34.4, 24. , 29. , 27.4, 22.8, 24.2, 29.7, 36.8,
       27.3, 28.2, 30.9, 37.5, 25.8, 37.8, 22.4, 48.9, 26.6, 27.2, 23.5,
       28.3, 44.2, 25.4, 22.2, 30.5, 26.5, 33.7, 23.1, 32. , 29.9, 23.9,
       28.5, 26.4, 20.2, 33.6, 38.6, 39.2, 27.7, 31.4, 36.5, 33.2, 32.8,
       40.4, 25.3, 30.2, 47.5, 20.3, 30. , 28.9, 28.1, 31.1, 21.7, 27. ,
       24.1, 45.9, 44.1, 22.9, 29.1, 32.3, 41.1, 25.6, 29.8, 26.3, 26.2,
       29.4, 24.4, 28. , 28.8, 34.6, 19.4, 30.3, 41.5, 22.6, 56.6, 27.1,
       31.3, 31. , 31.7, 35.8, 28.4, 20.1, 26.7, 38.7, 34.9, 25. , 23.8,
       21.8, 27.5, 24.6, 32.9, 26.1, 31.9, 34.1, 36.9, 37.3, 45.7, 34.2,
       23.6, 22.3, 37.1, 45. , 25.5, 30.8, 37.4, 34.5, 27.9, 29.5, 46. ,
       42.5, 35.5, 26.9, 45.5, 31.5, 33. , 23.4, 30.7, 20.5, 21.5, 40. ,
       28.6, 42.2, 29.6, 35.4, 16.9, 26.8, 39.3, 32.6, 35.9, 21.2, 42.4,
       40.5, 36.7, 29.3, 19.6, 18. , 17.6, 19.1, 50.1, 17.7, 54.6, 35. ,
       22. , 39.4, 19.7, 22.5, 25.2, 41.8, 60.9, 23.7, 24.5, 31.2, 16. ,
       31.6, 25.1, 24.8, 18.3, 20. , 19.5, 36. , 35.3, 40.1, 43.1, 21.4,
       34.3, 27.6, 16.5, 24.3, 25.7, 21.9, 38.4, 25.9, 54.7, 18.6, 24.9,
       48.2, 20.7, 39.5, 23.3, 64.8, 35.1, 43.6, 21. , 47.3, 16.6, 21.6,
       15.5, 35.6, 16.7, 41.9, 16.4, 17.1, 29.2, 37.9, 44.6, 39.6, 40.3,
       41.6, 39. , 23.2, 18.9, 36.1, 36.3, 46.5, 16.8, 46.6, 35.2, 20.9,
       13.8, 31.8, 15.3, 38.2, 45.2, 17. , 49.8, 27.8, 60.2, 23. , 22.1,
       26. , 44.3, 51. , 39.7, 34.7, 21.3, 41.2, 34.8, 19.2, 35.7, 40.8,
       24.7, 19. , 32.4, 34. , 28.7, 32.1, 51.5, 20.4, 30.6, 71.9, 19.3,
       40.9, 17.2, 16.1, 16.2, 40.6, 18.4, 21.1, 42.3, 32.2, 50.2, 17.5,
       18.7, 42.1, 47.8, 20.8, 30.1, 17.3, 36.4, 12. , 36.2, 55.7, 14.4,
       43. , 41.7, 33.8, 43.9, 22.7, 57.5, 37. , 38.5, 16.3, 44. , 32.7,
       54.2, 40.2, 33.3, 17.4, 41.3, 52.3, 14.6, 17.8, 46.1, 33.1, 18.1,
       43.8, 50.3, 38.9, 43.7, 39.9, 15.9, 19.8, 12.3, 78. , 38.3, 41. ,
       42.6, 43.4, 15.1, 20.6, 33.5, 43.2, 30.4, 38. , 33.4, 44.9, 44.7,
       37.6, 39.8, 53.4, 55.2, 42. , 37.2, 42.8, 18.8, 42.9, 14.3, 37.7,
       48.4, 50.6, 46.2, 49.5, 43.3, 33.9, 18.5, 44.5, 45.4, 55. , 54.8,
       19.9, 17.9, 15.6, 52.8, 15.2, 66.8, 55.1, 18.2, 48.5, 55.9, 57.3,
       10.3, 14.1, 15.7, 56. , 44.8, 13.4, 51.8, 38.1, 57.7, 44.4, 38.8,
       49.3, 39.1, 54. , 56.1, 97.6, 53.9, 13.7, 11.5, 41.4, 14.2, 49.4,
       15.4, 45.1, 49.2, 48.7, 53.8, 42.7, 48.8, 52.7, 53.5, 50.5, 15.8,
       45.3, 14.8, 51.9, 63.3, 40.7, 61.2, 48. , 46.8, 48.3, 58.1, 50.4,
       11.3, 12.8, 13.5, 14.5, 15. , 59.7, 47.4, 52.5, 13.2, 52.9, 61.6,
       49.9, 54.3, 47.9, 13. , 13.9, 50.9, 57.2, 64.4, 92. , 50.8, 57.9,
       45.8, 47.6, 14. , 46.4, 46.9, 47.1, 13.3, 48.1, 51.7, 46.3, 54.1,
       14.9])
```


```
df['bmi'].value_counts()
```




	count
bmi	
28.7	41
28.4	38
26.7	37
27.6	37
26.1	37
...	...
48.7	1
49.2	1
51.0	1
49.4	1
14.9	1

418 rows × 1 columns

```
df['bmi'].mean()
```

 28.893236911794666

```
df['bmi'].fillna(df['bmi'].mean(), inplace=True)
```

 <ipython-input-11-2a6d4795ba2f>:1: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment. The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col]

```
df['bmi'].fillna(df['bmi'].mean(), inplace=True)
```

```
df['bmi'].isnull().sum()
```


 0

```
type('bmi')
```

 str

```
df['bmi'] = pd.to_numeric(df['bmi'], errors='coerce')
```

```
df.info()
```


 <class 'pandas.core.frame.DataFrame'>  
RangeIndex: 5110 entries, 0 to 5109  
Data columns (total 12 columns):

#	Column	Non-Null Count	Dtype
0	id	5110 non-null	int64
1	gender	5110 non-null	object
2	age	5110 non-null	float64
3	hypertension	5110 non-null	int64
4	heart_disease	5110 non-null	int64
5	ever_married	5110 non-null	object
6	work_type	5110 non-null	object
7	Residence_type	5110 non-null	object
8	avg_glucose_level	5110 non-null	float64
9	bmi	5110 non-null	float64
10	smoking_status	5110 non-null	object
11	stroke	5110 non-null	int64

dtypes: float64(3), int64(4), object(5)  
memory usage: 479.2+ KB

## 1. Gender

```
df['gender'].unique()
```

 array(['Male', 'Female', 'Other'], dtype=object)

```
df['gender'].value_counts()
```

	count
gender	
Female	2994
Male	2115
Other	1

As there is only 1 value present for the OTHER, we will drop that

```
df = df.drop(df[df['gender'] == 'Other'].index)
```

```
df.head(5)
```

	id	gender	age	hypertension	heart_disease	ever_married	work_type	Residence_type	avg_glucose_level	bmi	smoking_s
0	9046	Male	67.0	0	1	Yes	Private	Urban	228.69	36.600000	formerly sr
1	51676	Female	61.0	0	0	Yes	Self-employed	Rural	202.21	28.893237	never sr
2	31112	Male	80.0	0	1	Yes	Private	Rural	105.92	32.500000	never sr
3	60182	Female	49.0	0	0	Yes	Private	Urban	171.23	34.400000	sr
4	1665	Female	79.0	1	0	Yes	Self-employed	Rural	174.12	24.000000	never sr

Next steps:

Generate code with df

View recommended plots

New interactive sheet

```
df['gender'].value_counts()
```

	count
gender	
Female	2994
Male	2115

Now replacing the Male as 1 and Female as 0

```
df['gender'].replace({'Male': 1, 'Female': 0}, inplace = True)
```

<ipython-input-21-a05e76498d29>:1: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained ass  
The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col]

```
df['gender'].replace({'Male': 1, 'Female': 0}, inplace = True)  
<ipython-input-21-a05e76498d29>:1: FutureWarning: Downcasting behavior in `replace` is deprecated and will be removed in a future ve  
df['gender'].replace({'Male': 1, 'Female': 0}, inplace = True)
```

```
df.head(3)
```

	id	gender	age	hypertension	heart_disease	ever_married	work_type	Residence_type	avg_glucose_level	bmi	smoking_s
0	9046	1	67.0	0	1	Yes	Private	Urban	228.69	36.600000	formerly sr
1	51676	0	61.0	0	0	Yes	Self-employed	Rural	202.21	28.893237	never sr
2	31112	1	80.0	0	1	Yes	Private	Rural	105.92	32.500000	never sr

Next steps:

Generate code with df

View recommended plots

New interactive sheet

## 2. Ever\_married

```
df['ever_married'].unique()
```

```
array(['Yes', 'No'], dtype=object)
```

```
df['ever_married'].value_counts()
```

```
count
ever_married
Yes      3353
No       1756
```

### Replacing Yes as 1 and No and 0

```
df['ever_married'].replace({'Yes' : 1, 'No' : 0}, inplace = True)
```

```
<ipython-input-25-8391b8ec88f2>:1: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment. The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting
```

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col]

```
df['ever_married'].replace({'Yes' : 1, 'No' : 0}, inplace = True)
<ipython-input-25-8391b8ec88f2>:1: FutureWarning: Downcasting behavior in `replace` is deprecated and will be removed in a future version.
df['ever_married'].replace({'Yes' : 1, 'No' : 0}, inplace = True)
```

```
df.head(3)
```

```
id  gender  age  hypertension  heart_disease  ever_married  work_type  Residence_type  avg_glucose_level  bmi  smoking_status
0   9046    1  67.0            0             1           1   Private      Urban              228.69  36.600000  formerly sr
1   51676   0  61.0            0             0           1  Self-employed  Rural              202.21  28.893237  never sr
2   31112   1  80.0            0             1           1   Private      Rural              105.92  32.500000  never sr
```

Next steps:

[Generate code with df](#)[View recommended plots](#)[New interactive sheet](#)

## 3. Work Type

```
df['work_type'].unique()
```

```
array(['Private', 'Self-employed', 'Govt_job', 'children', 'Never_worked'],
      dtype=object)
```

```
df['work_type'].value_counts()
```

```
count
work_type
Private      2924
Self-employed  819
children      687
Govt_job      657
Never_worked   22
```

## 4. Residence type

```
df['Residence_type'].unique()
```

```
array(['Urban', 'Rural'], dtype=object)
```

```
df['Residence_type'].value_counts()
```

```
count
Residence_type
Urban      2596
Rural      2513
```

```
df['Residence_type'].replace({'Urban' : 1, 'Rural' : 0}, inplace = True)
```

```
<ipython-input-31-fadd75f3552d>:1: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment. The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col]
```

```
df['Residence_type'].replace({'Urban' : 1, 'Rural' : 0}, inplace = True)
<ipython-input-31-fadd75f3552d>:1: FutureWarning: Downcasting behavior in `replace` is deprecated and will be removed in a future version. df['Residence_type'].replace({'Urban' : 1, 'Rural' : 0}, inplace = True)
```

```
df.head(3)
```

```
id gender age hypertension heart_disease ever_married work_type Residence_type avg_glucose_level bmi smoking_status
0  9046    1  67.0          0             1           1   Private              1      228.69  36.600000  formerly smoked
1  51676    0  61.0          0             0           1  Self-employed              0      202.21  28.893237  never smoked
2  31112    1  80.0          0             1           1   Private              0      105.92  32.500000  never smoked
```

Next steps: [Generate code with df](#) [View recommended plots](#) [New interactive sheet](#)

## 5. Smoking status

```
df['smoking_status'].unique()
```

```
array(['formerly smoked', 'never smoked', 'smokes', 'Unknown'],
      dtype=object)
```

```
df['smoking_status'].value_counts()
```

```
count
smoking_status
never smoked    1892
Unknown         1544
formerly smoked   884
smokes          789
```

## 6. ID

I will drop this feature as it has no impact on my target feature.

```
df.drop('id', axis = 1, inplace = True)
```

```
df.head(3)
```

	gender	age	hypertension	heart_disease	ever_married	work_type	Residence_type	avg_glucose_level	bmi	smoking_status	:
0	1	67.0	0	1	1	Private	1	228.69	36.600000	formerly smoked	
1	0	61.0	0	0	1	Self-employed	0	202.21	28.893237	never smoked	

Next steps:

[Generate code with df](#)

[View recommended plots](#)

[New interactive sheet](#)

df.dtypes

		0
gender		int64
age		float64
hypertension		int64
heart_disease		int64
ever_married		int64
work_type		object
Residence_type		int64
avg_glucose_level		float64
bmi		float64
smoking_status		object
stroke		int64

Target Column (Stroke) Our target column consist of 0 and 1 which result to Binary Classification Problem

df['stroke'].value\_counts()


	count
stroke	
0	4860
1	249

df['stroke'].value\_counts(normalize = True) \* 100

	proportion
stroke	
0	95.126248
1	4.873752


The dataset is in imbalance form as the proportion of not having stroke is 95% whereas the proportion of having stroke is only 5%. Hence we need to fix this.

```
stroke_counts = df['stroke'].value_counts()
sns.barplot(x=stroke_counts.index, y=stroke_counts.values)
```




We will use ROC metric to fix the imbalance target column dataset.

```
sns.histplot(x = df['age'], kde = True)
```



```
gender_counts = df['gender'].value_counts()  
sns.barplot(x=gender_counts.index, y=gender_counts.values)
```





Male are having higher chances of getting stroke as compared to female

Dividing the data as per categorical and numerical features

```
categorical_features = [x for x in df.columns if df[x].dtypes == 'O']
```

```
categorical_features
```

```
['work_type', 'smoking_status']
```

```
numerical_features = [x for x in df.columns if df[x].dtypes != 'O']
```

```
numerical_features
```

```
['gender',
 'age',
 'hypertension',
 'heart_disease',
 'ever_married',
 'Residence_type',
 'avg_glucose_level',
 'bmi',
 'stroke']
```

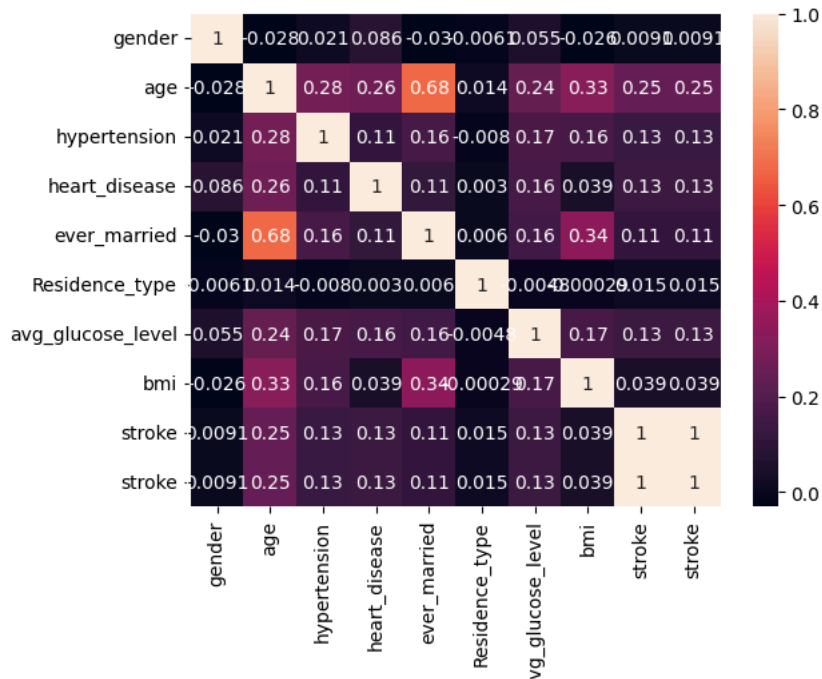
Checking the correlation between the dependent and independent variables

```
df[numerical_features + ['stroke']].corr()
```

	gender	age	hypertension	heart_disease	ever_married	Residence_type	avg_glucose_level	bmi	stroke
gender	1.000000	-0.027752	0.021223	0.085685	-0.030171	-0.006105	0.054722	-0.025606	0.0090
age	-0.027752	1.000000	0.276367	0.263777	0.679084	0.014031	0.238323	0.325858	0.2452
hypertension	0.021223	0.276367	1.000000	0.108292	0.164187	-0.007980	0.174540	0.160147	0.1278
heart_disease	0.085685	0.263777	0.108292	1.000000	0.114601	0.003045	0.161907	0.038862	0.1349
ever_married	-0.030171	0.679084	0.164187	0.114601	1.000000	0.005988	0.155329	0.335563	0.1082
Residence_type	-0.006105	0.014031	-0.007980	0.003045	0.005988	1.000000	-0.004783	-0.000288	0.0154
avg_glucose_level	0.054722	0.238323	0.174540	0.161907	0.155329	-0.004783	1.000000	0.168910	0.1319
bmi	-0.025606	0.325858	0.160147	0.038862	0.335563	-0.000288	0.168910	1.000000	0.0389
stroke	0.009081	0.245239	0.127891	0.134905	0.108299	0.015415	0.131991	0.038912	1.0000

```
sns.heatmap(df[numerical_features + ['stroke']].corr(), annot = True)
```

<Axes: >



There is no features which are having higher correlation with our target column.

numerical\_features

```
['gender',  
 'age',  
 'hypertension',  
 'heart_disease',  
 'ever_married',  
 'Residence_type',  
 'avg_glucose_level',  
 'bmi',  
 'stroke']
```

categorical\_features

```
['work_type', 'smoking_status']
```

df.head(3)

```
gender  age  hypertension  heart_disease  ever_married  work_type  Residence_type  avg_glucose_level  bmi  smoking_status  
0      1    67.0           0              1              1    Private           228.69    36.600000    formerly smoked  
1      0    61.0           0              0              1    Self-employed       202.21    28.893237    never smoked
```

Next steps:

[Generate code with df](#)

[View recommended plots](#)

[New interactive sheet](#)

df['bmi'].shape

```
(5109,)
```

Converting my categorical columns to numerical columns

categorical\_features

```
['work_type', 'smoking_status']
```

df.info()

```
<class 'pandas.core.frame.DataFrame'>  
Index: 5109 entries, 0 to 5109  
Data columns (total 11 columns):  
#   Column      Non-Null Count  Dtype  
---  -  
0   gender      5109 non-null   int64
```

```

1  age          5109 non-null  float64
2  hypertension  5109 non-null  int64
3  heart_disease 5109 non-null  int64
4  ever_married  5109 non-null  int64
5  work_type     5109 non-null  object
6  Residence_type 5109 non-null  int64
7  avg_glucose_level 5109 non-null float64
8  bmi          5109 non-null  float64
9  smoking_status 5109 non-null  object
10 stroke       5109 non-null  int64
dtypes: float64(3), int64(6), object(2)
memory usage: 479.0+ KB

```

```
df.isnull().sum()
```

```

gender    0
age       0
hypertension  0
heart_disease  0
ever_married  0
work_type  0
Residence_type  0
avg_glucose_level  0
bmi          0
smoking_status  0
stroke       0

```

```

# Drop rows with any null values in the relevant columns
df = df.dropna(subset=['work_type', 'smoking_status'])

```

```

# Ensure proper data types
df['work_type'] = df['work_type'].astype(str)
df['smoking_status'] = df['smoking_status'].astype(str)

```

```

# Create a OneHotEncoder instance
ohe = OneHotEncoder()

```

```

# Fit and transform the categorical columns
encoded_features = ohe.fit_transform(df[['work_type', 'smoking_status']])

```

```

# Create a DataFrame with the new encoded features
encoded_df = pd.DataFrame(encoded_features.toarray(), columns=ohe.get_feature_names_out(['work_type', 'smoking_status']))

```

```

# Concatenate the original DataFrame with the new one, dropping the original columns
df = pd.concat([df.drop(['work_type', 'smoking_status'], axis=1), encoded_df], axis=1)

```

```
print(df)
```

```

gender  age  hypertension  heart_disease  ever_married  Residence_type  \
0      1.0  67.0          0.0           1.0           1.0           1.0
1      0.0  61.0          0.0           0.0           1.0           0.0
2      1.0  80.0          0.0           1.0           1.0           0.0
3      0.0  49.0          0.0           0.0           1.0           1.0
4      0.0  79.0          1.0           0.0           1.0           0.0
...     ...     ...          ...          ...          ...          ...
5106    0.0  81.0          0.0           0.0           1.0           1.0
5107    0.0  35.0          0.0           0.0           1.0           0.0
5108    1.0  51.0          0.0           0.0           1.0           0.0
5109    0.0  44.0          0.0           0.0           1.0           1.0
3116   NaN   NaN          NaN           NaN           NaN           NaN

avg_glucose_level  bmi  stroke  work_type_Govt_job  \
0      228.69  36.600000  1.0      0.0
1      202.21  28.893237  1.0      0.0
2      105.92  32.500000  1.0      0.0
3      171.23  34.400000  1.0      0.0
4      174.12  24.000000  1.0      0.0
...     ...     ...          ...          ...
5106    125.20  40.000000  0.0      0.0
5107     82.99  30.600000  0.0      0.0
5108    166.29  25.600000  0.0      1.0
5109     85.28  26.200000  0.0     NaN
3116      NaN      NaN      NaN      0.0

```

	work_type_Never_worked	work_type_Private	work_type_Self-employed	\
0	0.0	1.0	0.0	
1	0.0	0.0	1.0	
2	0.0	1.0	0.0	
3	0.0	1.0	0.0	
4	0.0	0.0	1.0	
...	...	...	...	
5106	0.0	0.0	1.0	
5107	0.0	1.0	0.0	
5108	0.0	0.0	0.0	
5109	NaN	NaN	NaN	
3116	0.0	0.0	0.0	

	work_type_children	smoking_status_Unknown	\
0	0.0	0.0	
1	0.0	0.0	
2	0.0	0.0	
3	0.0	0.0	
4	0.0	0.0	
...	...	...	
5106	0.0	0.0	
5107	0.0	0.0	
5108	0.0	1.0	
5109	NaN	NaN	
3116	1.0	1.0	

	smoking_status_formerly smoked	smoking_status_never smoked	\
0	1.0	0.0	
1	0.0	1.0	
2	0.0	1.0	
3	0.0	0.0	
4	0.0	1.0	

```
df = df.drop(df.index[-2:])
```

```
df.isnull().sum()
```

	0
gender	0
age	0
hypertension	0
heart_disease	0
ever_married	0
Residence_type	0
avg_glucose_level	0
bmi	0
stroke	0
work_type_Govt_job	0
work_type_Never_worked	0
work_type_Private	0
work_type_Self-employed	0
work_type_children	0
smoking_status_Unknown	0
smoking_status_formerly smoked	0
smoking_status_never smoked	0
smoking_status_smokes	0


```
df.head(3)
```

	gender	age	hypertension	heart_disease	ever_married	Residence_type	avg_glucose_level	bmi	stroke	work_type_Govt_job
0	1.0	67.0	0.0	1.0	1.0	1.0	228.69	36.600000	1.0	0.0
1	0.0	61.0	0.0	0.0	1.0	0.0	202.21	28.893237	1.0	0.0
2	1.0	80.0	0.0	1.0	1.0	0.0	105.92	32.500000	1.0	0.0

Next steps:

[Generate code with df](#)[View recommended plots](#)[New interactive sheet](#)

```
sns.heatmap(df[numerical_features].corr(), annot = True)
```



There is no correlation hence we can go for model prediction

```
# Define features and target
X = df.drop('stroke', axis=1)
y = df['stroke']

# Split data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

As our target column is imabalnce we are fixing this using SMOTE

```
# Apply SMOTE to the training data
smote = SMOTE(random_state=42)
X_train_resampled, y_train_resampled = smote.fit_resample(X_train, y_train)

# Check the class distribution after SMOTE
print("Original training set shape:", X_train.shape, "Class distribution:", y_train.value_counts())
print("Resampled training set shape:", X_train_resampled.shape, "Class distribution:", pd.Series(y_train_resampled).value_counts())
```

```
Original training set shape: (4086, 17) Class distribution: stroke
0.0    3899
1.0     187
Name: count, dtype: int64
Resampled training set shape: (7798, 17) Class distribution: stroke
0.0    3899
1.0    3899
Name: count, dtype: int64
```

```
df.head(3)
```

<Axes: >

	gender	age	hypertension	heart_disease	ever_married	Residence_type	avg_glucose_level	bmi	stroke	work_type_Govt_job
0	1.0	67.0	0.0	1.0	1.0	1.0	228.69	36.600000	1.0	0.0
1	0.0	61.0	0.0	0.0	1.0	0.0	202.21	28.893237	1.0	0.0
2	1.0	80.0	0.0	1.0	1.0	0.0	105.92	32.500000	1.0	0.0

Next steps:

[Generate code with df](#)[View recommended plots](#)[New interactive sheet](#)

```
# Scale features
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

# Train model
model = RandomForestClassifier(random_state=42)
model.fit(X_train, y_train)
```

```
RandomForestClassifier
RandomForestClassifier(random_state=42)
```

```
# Make predictions
y_pred = model.predict(X_test)
```

```
# Evaluate model
print(confusion_matrix(y_test, y_pred))
print(classification_report(y_test, y_pred))
```

```
[[958  2]
 [ 62  0]]
```

	precision	recall	f1-score	support
0.0	0.94	1.00	0.97	960
1.0	0.00	0.00	0.00	62
accuracy			0.94	1022
macro avg	0.47	0.50	0.48	1022
weighted avg	0.88	0.94	0.91	1022

## Implementing KNN

```
# List to store the accuracy for different values of k
k_values = range(1, 21) # Test k from 1 to 20
accuracy = []

# Loop to find the ideal k
for k in k_values:
    knn = KNeighborsClassifier(n_neighbors=k)
    # Perform 5-fold cross-validation and get the mean accuracy
    cv_scores = cross_val_score(knn, X_train, y_train, cv=5)
    accuracy.append(cv_scores.mean())
```

```
# Plot the results
plt.figure(figsize=(10, 6))
plt.plot(k_values, accuracy, marker='o')
plt.title('KNN: Finding the Ideal k')
plt.xlabel('Number of Neighbors (k)')
plt.ylabel('Mean Accuracy')
plt.xticks(k_values)
plt.grid()
plt.show()
```

